

VIBOT 2018

Closing remarks

Francesco Ciompi

francesco.ciompi@radboudumc.nl

VGG-net (aka OxfordNet), 2014

http://www.robots.ox.ac.uk/~vgg/research/very_deep/

VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION

Karen Simonyan* & Andrew Zisserman⁺

Visual Geometry Group, Department of Engineering Science, University of Oxford
`{karen, az}@robots.ox.ac.uk`

- 19 layers
- **Fixed filter size: 3x3**

ABSTRACT

In this work we investigate the effect of the convolutional network depth on its accuracy in the large-scale image recognition setting. Our main contribution is a thorough evaluation of networks of increasing depth using an architecture with very small (3×3) convolution filters, which shows that a significant improvement on the prior-art configurations can be achieved by pushing the depth to 16–19 weight layers. These findings were the basis of our ImageNet Challenge 2014 submission, where our team secured the first and the second places in the localisation and classification tracks respectively. We also show that our representations generalise well to other datasets, where they achieve state-of-the-art results. We have made our two best-performing ConvNet models publicly available to facilitate further research on the use of deep visual representations in computer vision.

VGG-net

Results on ILSVRC-2014

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64	conv3-64	conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128	conv3-128	conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

ConvNet config. (Table 1)	smallest image side		top-1 val. error (%)	top-5 val. error (%)
	train (S)	test (Q)		
A	256	256	29.6	10.4
A-LRN	256	256	29.7	10.5
B	256	256	28.7	9.9
	256	256	28.1	9.4
C	384	384	28.1	9.3
	[256;512]	384	27.3	8.8
	256	256	27.0	8.8
D	384	384	26.8	8.7
	[256;512]	384	25.6	8.1
	256	256	27.3	9.0
E	384	384	26.9	8.7
	[256;512]	384	25.5	8.0

GoogLeNet, 2015

Going deeper with convolutions

Christian Szegedy

Google Inc.

Wei Liu

University of North Carolina, Chapel Hill

Yangqing Jia

Google Inc.

Pierre Sermanet

Google Inc.

Scott Reed

University of Michigan

Dragomir Anguelov

Google Inc.

Dumitru Erhan

Google Inc.

Vincent Vanhoucke

Google Inc.

Andrew Rabinovich

Google Inc.

Abstract

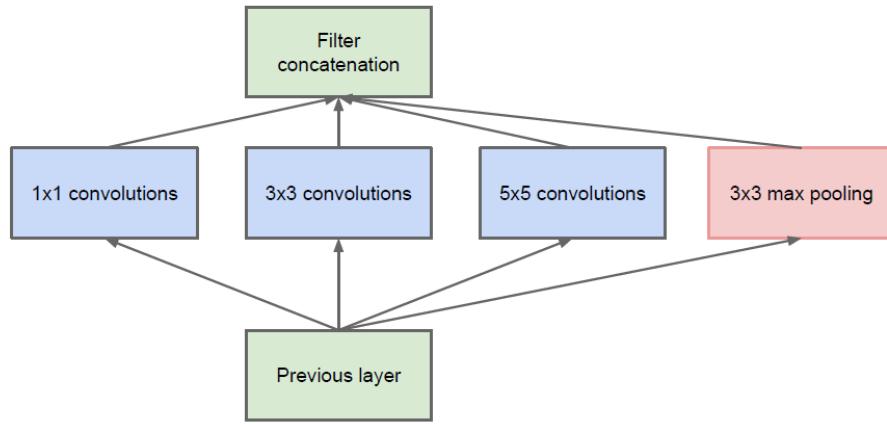
We propose a deep convolutional neural network architecture codenamed Inception, which was responsible for setting the new state of the art for classification and detection in the ImageNet Large-Scale Visual Recognition Challenge 2014 (ILSVRC14). The main hallmark of this architecture is the improved utilization of the computing resources inside the network. This was achieved by a carefully crafted design that allows for increasing the depth and width of the network while keeping the computational budget constant. To optimize quality, the architectural decisions were based on the Hebbian principle and the intuition of multi-scale processing. One particular incarnation used in our submission for ILSVRC14 is called GoogLeNet, a 22 layers deep network, the quality of which is assessed in the context of classification and detection.

- 22 layers
- Introduces a new layer architecture

GoogLeNet

This is the **new layer architecture**

The outputs are **concatenated**



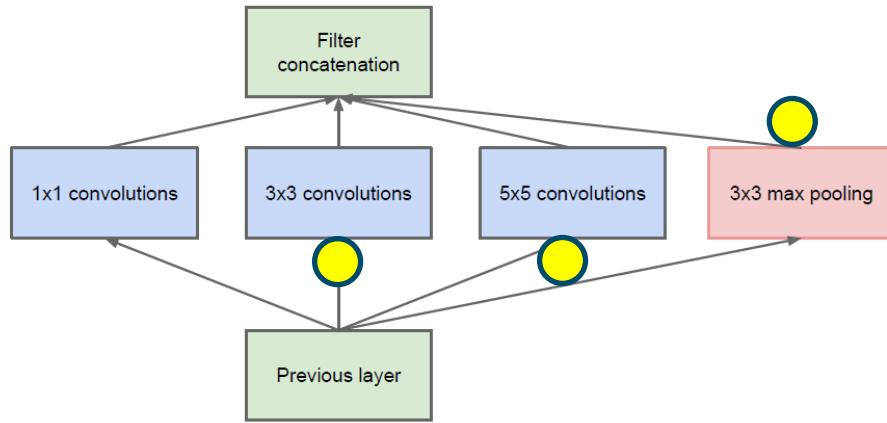
Problems

- Concatenation makes the matrix very big
- max-pooling does not change the number of feature maps, only their size
- **We should reduce the feature maps at some locations in the architecture**

GoogLeNet

This is the new layer architecture

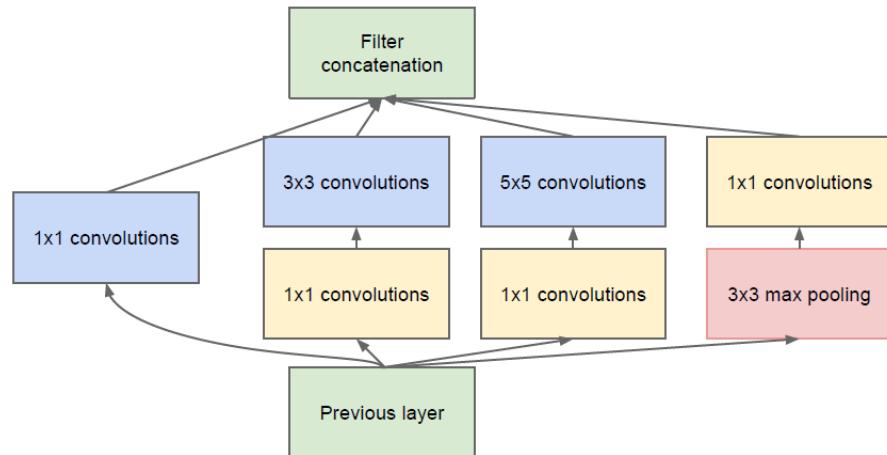
The outputs are **concatenated**



Problems

- Concatenation makes the matrix very big
- max-pooling does not change the number of feature maps, only their size
- **We should reduce the feature maps at some locations in the architecture**

GoogLeNet



The solution is using additional 1x1 convolutions!

It does not change the feature map size

We can tune the number of output feature maps: **dimension reduction**

It combines (+ non-linearity) all activations across feature maps and makes a smaller set

GoogLeNet

About 1×1 convolutions

- Given a feature map of $N \times N \times M$
- Given K $1 \times 1 \times M$ convolutional filters
- The result of the convolution between the maps and the filters has a size of:

$$N \times N \times K$$

GoogLeNet

About 1×1 convolutions

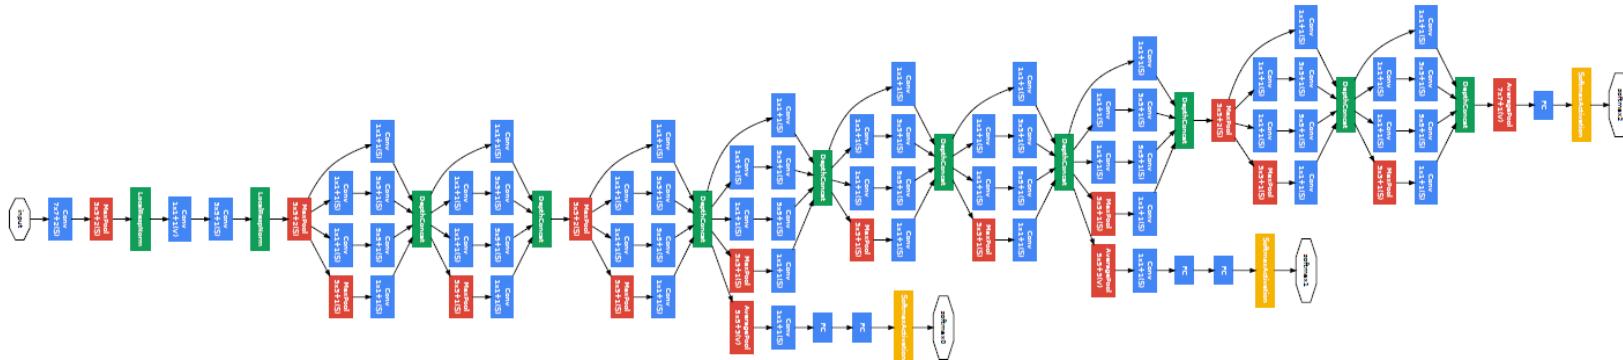
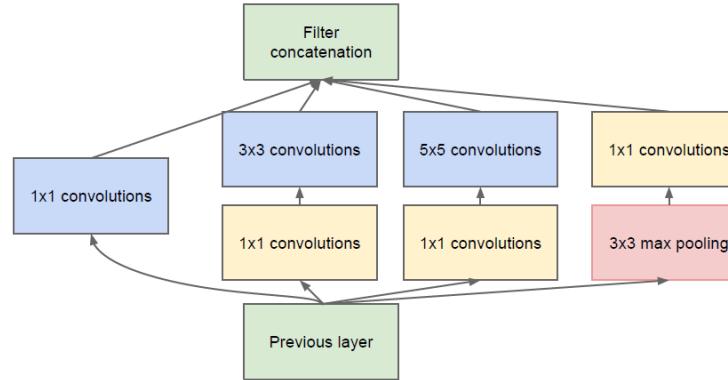
- Given a feature map of $N \times N \times M$
- Given K $1 \times 1 \times M$ convolutional filters
- The result of the convolution between the maps and the filters has a size of:

$N \times N \times K$

1x1 convolutions allow to change the “depth” of feature maps!

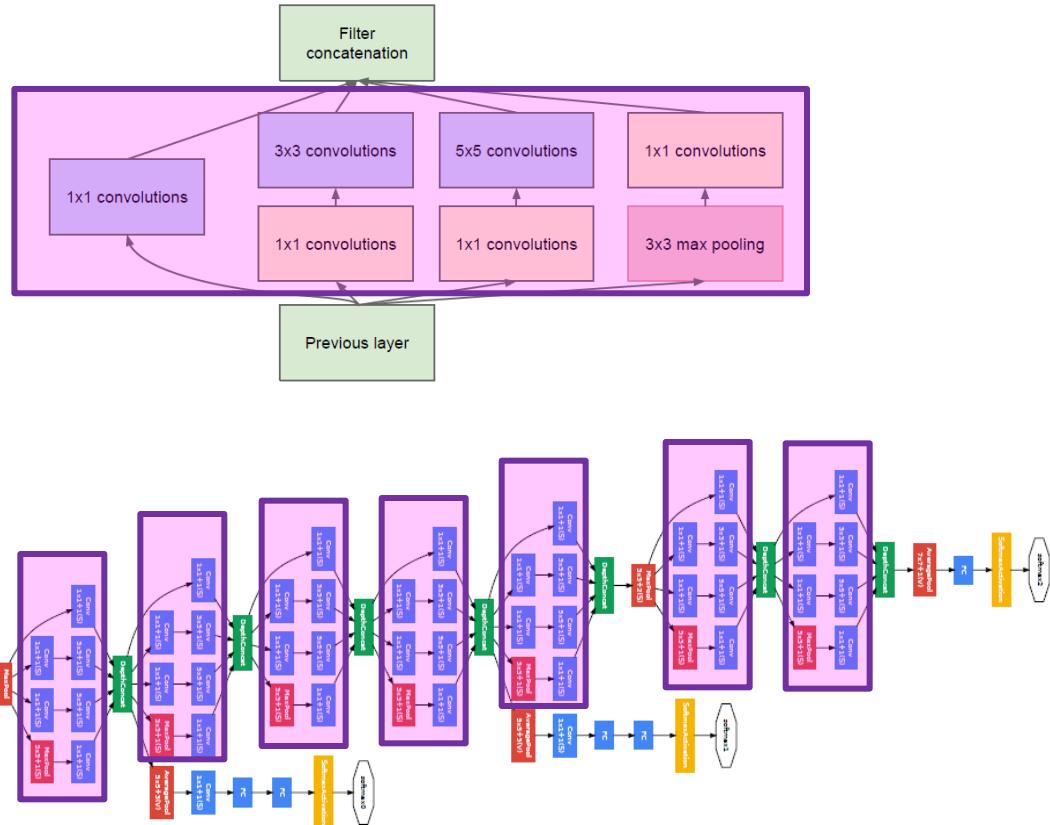
GoogLeNet

- Inception network



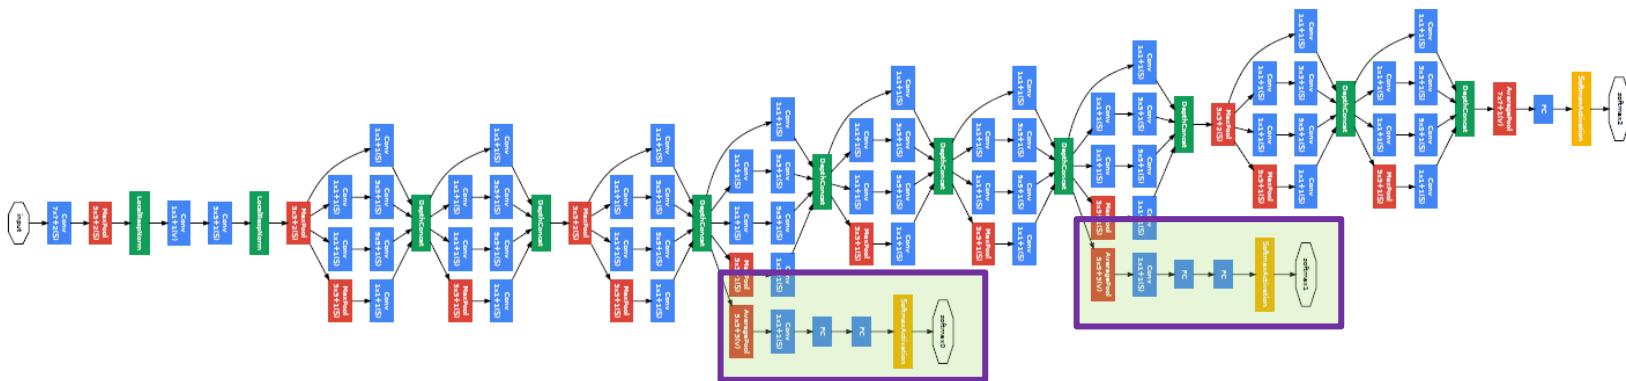
GoogLeNet

- Inception network



GoogLeNet

- Two additional classifiers during training
 - Since the network is very deep, it might not propagate gradients back through all the layers in an effective manner
- Additional classifiers:
 - encourage discrimination in the lower stages
 - increase the gradient signal that gets propagated back



GoogLeNet

- Results on ILSVRC-2014

Team	Year	Place	Error (top-5)	Uses external data
SuperVision	2012	1st	16.4%	no
SuperVision	2012	1st	15.3%	Imagenet 22k
Clarifai	2013	1st	11.7%	no
Clarifai	2013	1st	11.2%	Imagenet 22k
MSRA	2014	3rd	7.35%	no
VGG	2014	2nd	7.32%	no
GoogLeNet	2014	1st	6.67%	no

Table 2: Classification performance

architecture											
type	patch size/stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								

Table 1: GoogLeNet incarnation of the Inception architecture

GoogLeNet

- **Average pooling** replaces fully-connected layers

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								



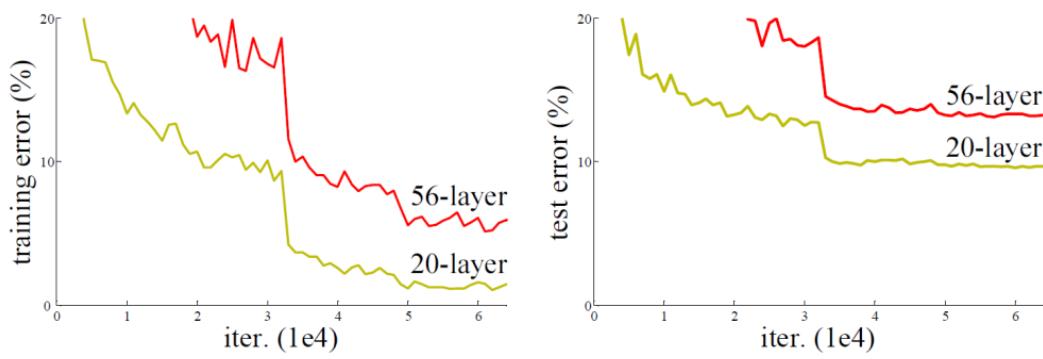


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

(He et al., 2015)

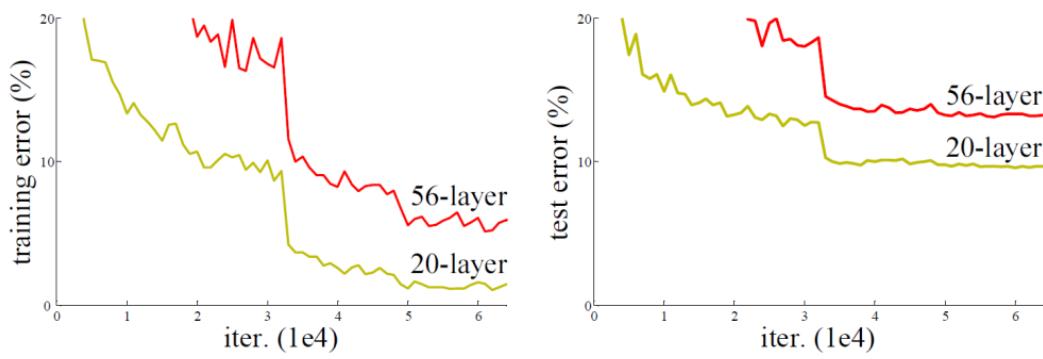


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

(He et al., 2015)



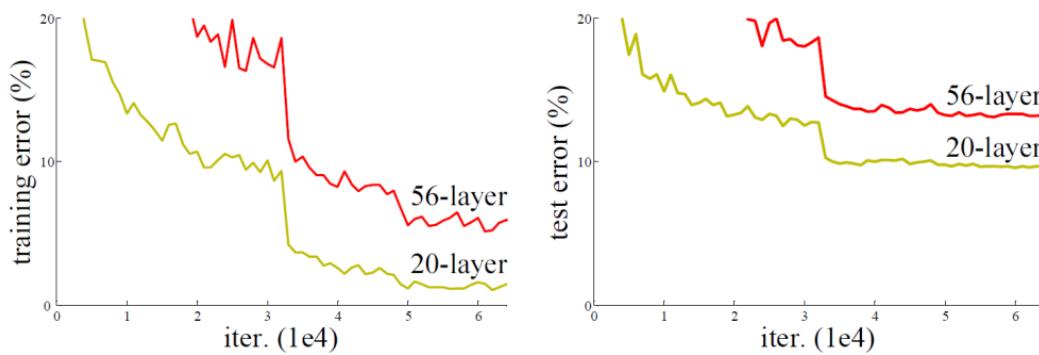
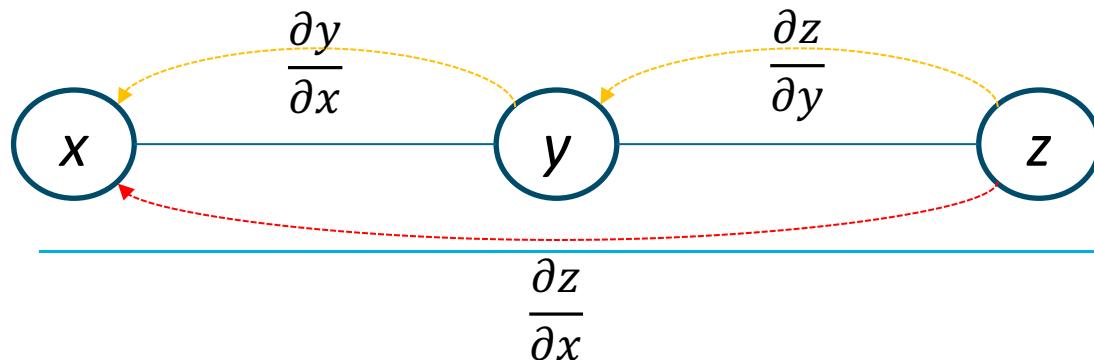


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

(He et al., 2015)



BACKPROPAGATION

ResNet, 2015

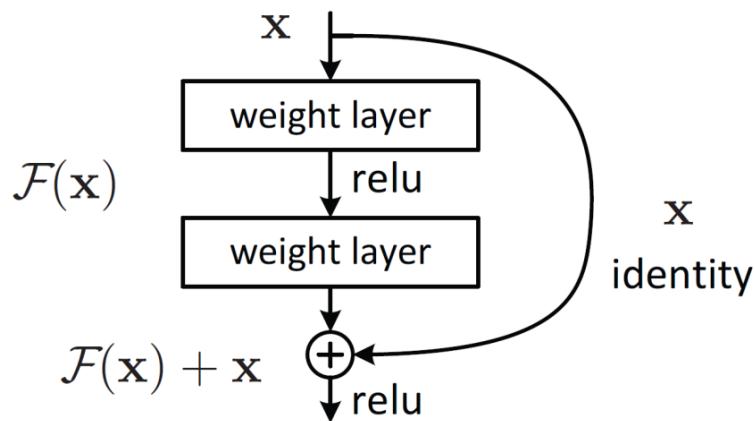
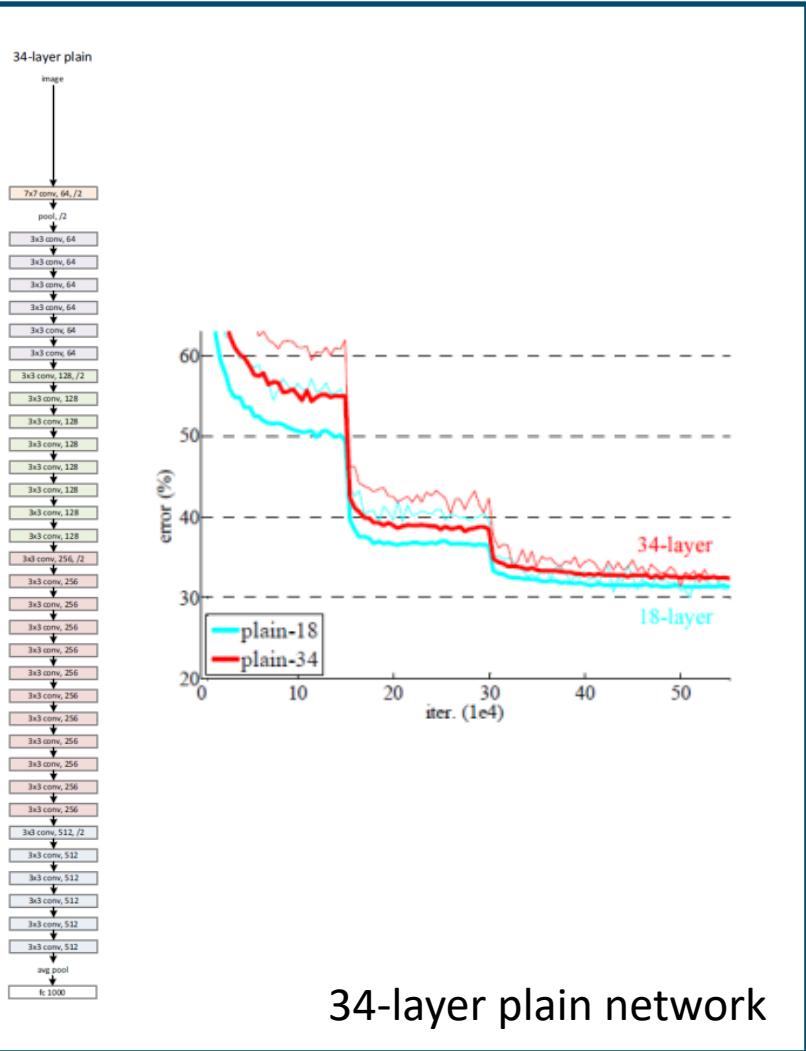


Figure 2. Residual learning: a building block.

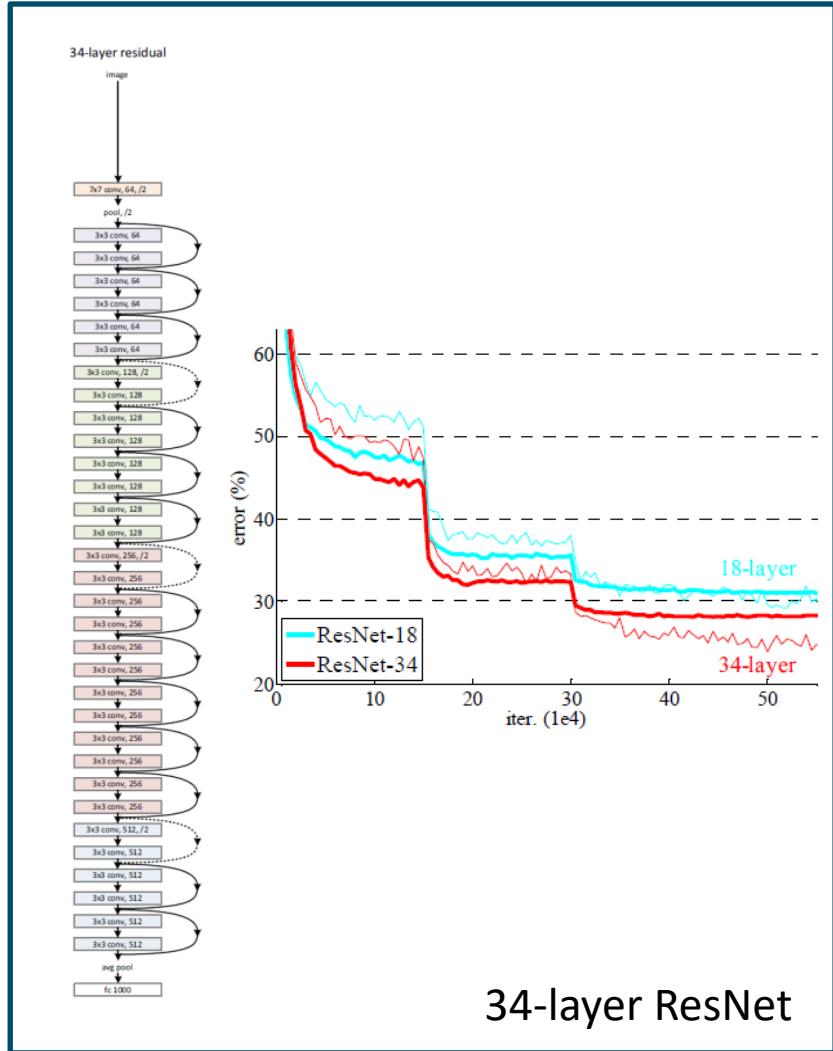
Gradients do not flow back to
the input signal
Skip connection helps!

As we go deep, representation
difficult to learn
Learning the residual helps!

ResNet

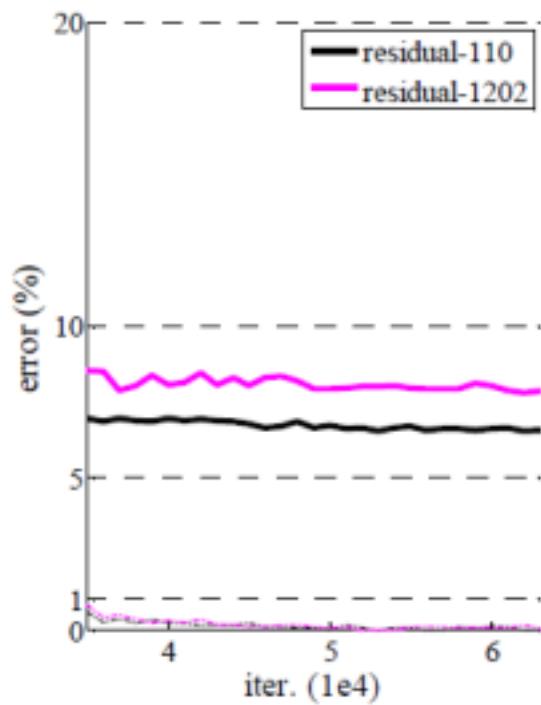
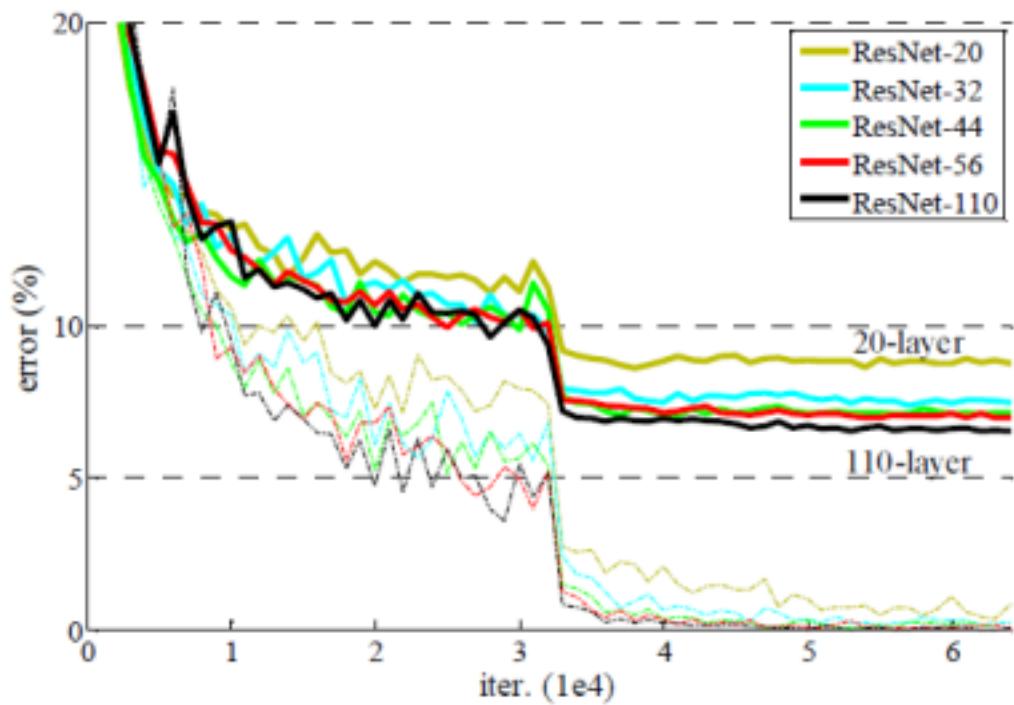


34-layer plain network

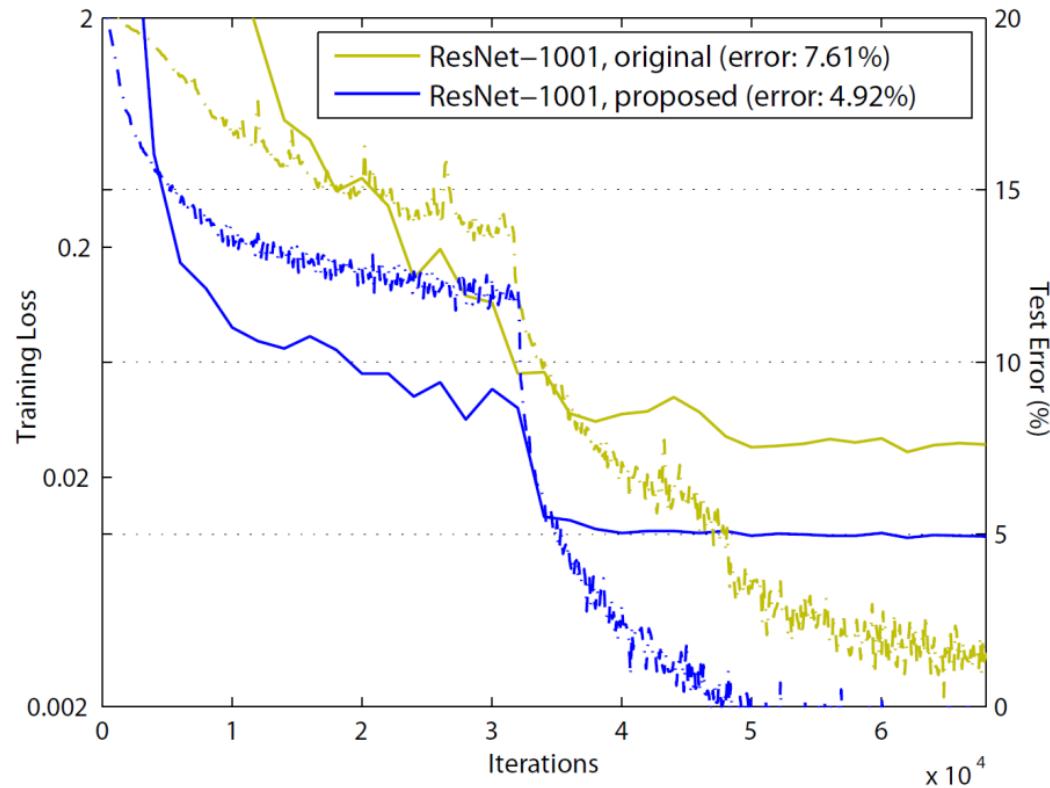
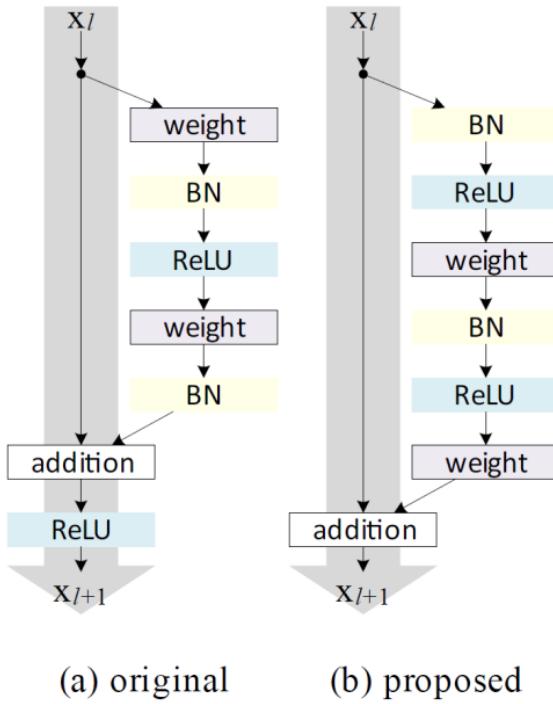


34-layer ResNet

ResNet



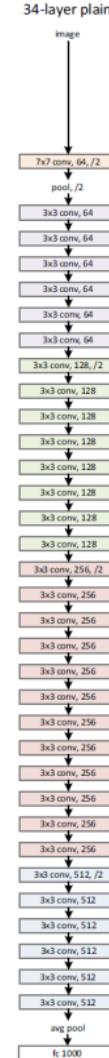
ResNet, 2016



(He et al., 2016)

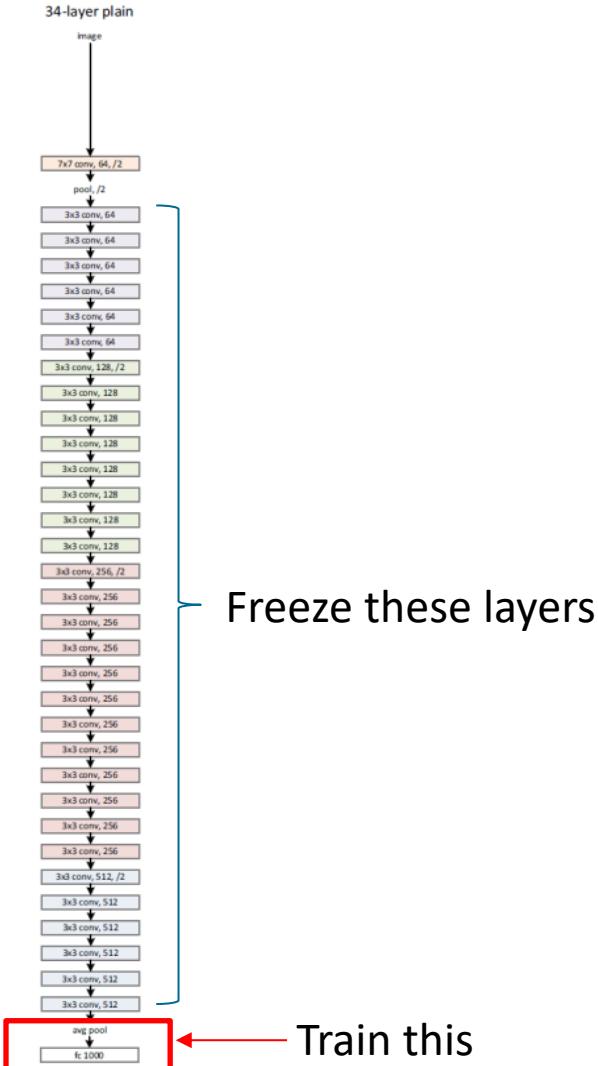
Using pre-trained networks

- VGG16, Inception, ResNet **trained on ImageNet** are available
- We can use them!
- This can help when we have limited training data



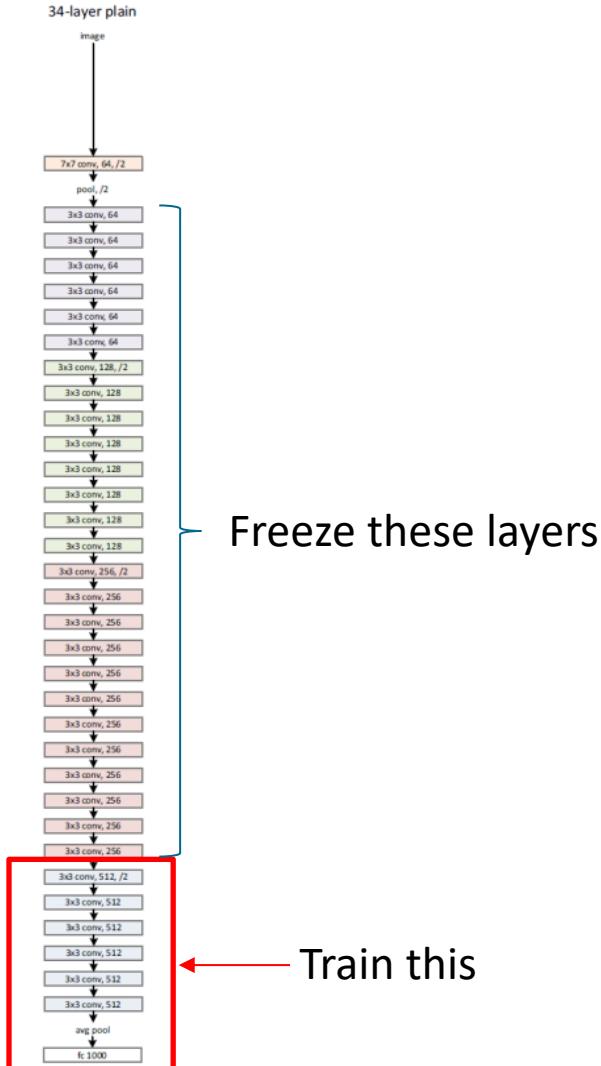
Using pre-trained networks

- VGG16, Inception, ResNet **trained on ImageNet** are available
- We can use them!
- This can help when we have limited training data
- Very **small** dataset:
 - **Feature extractor**



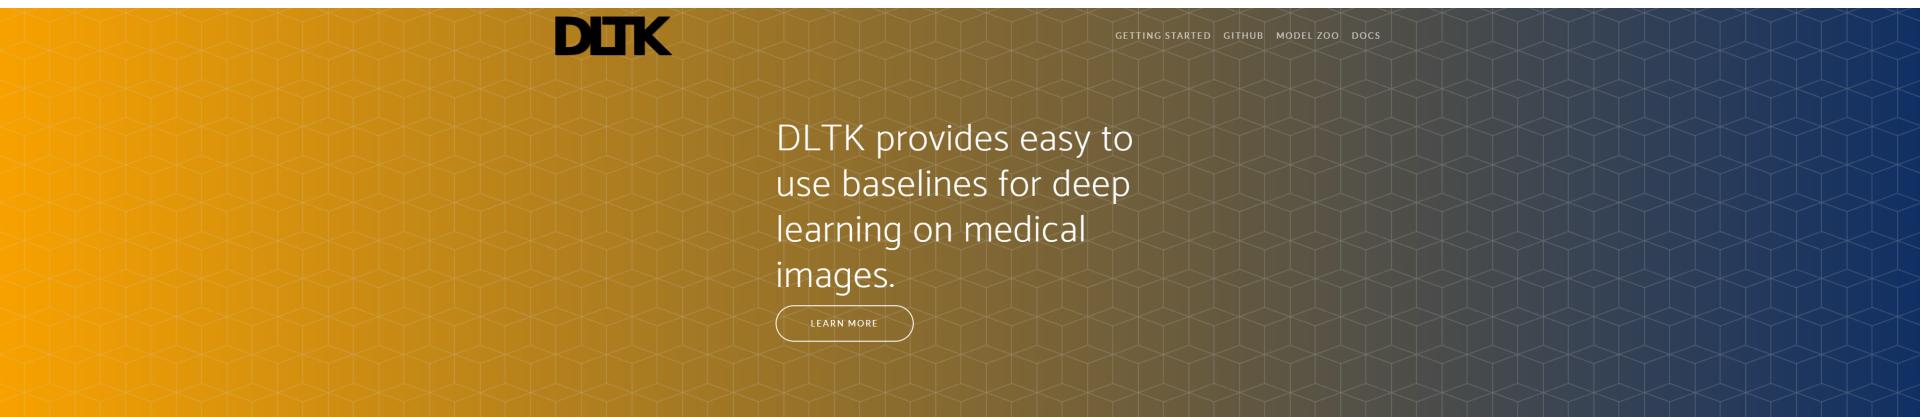
Using pre-trained networks

- VGG16, Inception, ResNet **trained on ImageNet** are available
- We can use them!
- This can help when we have limited training data
- Very small dataset:
 - Feature extractor
- **Medium** dataset:
 - **Fine-tuning**



Deep Learning in Medical Imaging

- Software packages



Getting Started.

DLTK is an open source library that makes deep learning on medical images easier. Built on TensorFlow, it enables fast prototyping and is simply installed via pip:

```
pip install dltk
```

DLTK comes with introduction [tutorials](#) and basic sample [applications](#), including scripts to download data.

Deep Learning in Medical Imaging

- Software packages

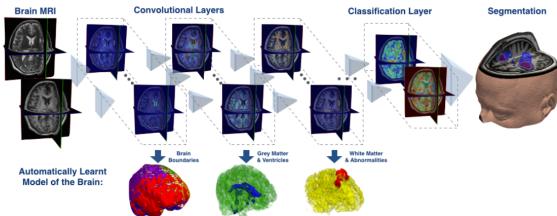
 BioMedIA About ▾ Research ▾ People Vacancies Internal

Home / Software / DeepMedic

DeepMedic

What's DeepMedic?

DeepMedic is our software for brain lesion segmentation based on a multi-scale 3D Deep Convolutional Neural Network coupled with a 3D fully connected Conditional Random Field. The system has been shown to yield excellent performance (winner of the ISLES 2015 competition) on challenging lesion segmentation tasks, including traumatic brain injuries, brain tumors, and ischemic stroke lesions.



Can I try it?

Yes! The software has been released open source on [Github](#).

Where can I find the details?

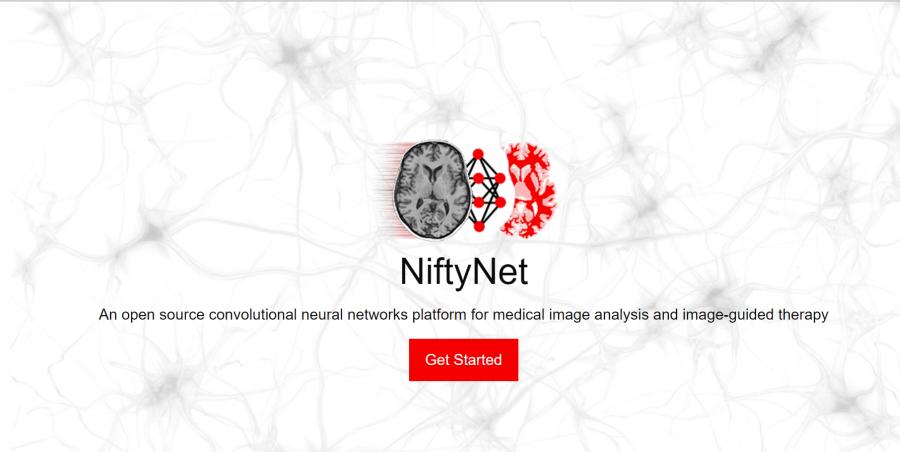
In our [open access article in Medical Image Analysis](#).

Who to ask?

Any questions? Get in touch with [Konstantinos Kamnitsas](#) or [Ben Glocker](#)

Deep Learning in Medical Imaging

- Software packages



What is NiftyNet?

NiftyNet is a [TensorFlow](#)-based open-source convolutional neural networks (CNNs) platform for research in medical image analysis and image-guided therapy. NiftyNet's modular structure is designed for sharing networks and pre-trained models. Using this modular structure you can:

- Get started with established pre-trained networks using built-in tools;
- Adapt existing networks to your imaging data;
- Quickly build new solutions to your own image analysis problems.

The code is available via [GitLab](#), or you can quickly get started with the [PyPI](#) module available [here](#).

[About this Specialization](#)[Courses](#)[Creators](#)[FAQ](#)[Enroll](#)

Started Apr 09

[Apply for Financial Aid](#)

Deep Learning Specialization

Master Deep Learning, and Break into AI

[About This Specialization](#)

CS231n: Convolutional Neural Networks for Visual Recognition

Spring 2018

Previous Years: [\[Winter 2015\]](#) [\[Winter 2016\]](#) [\[Spring 2017\]](#)



Course Description

Computer Vision has become ubiquitous in our society, with applications in search, image understanding, apps, mapping, medicine, drones, and self-driving cars. Core to many of these applications are visual recognition tasks such as image classification, localization and detection. Recent developments in neural network (aka "deep learning") approaches have greatly advanced the performance of these state-of-the-art visual recognition systems. This course is a deep dive into details of the deep learning architectures with a focus on learning end-to-end models for these tasks, particularly image classification. During the 10-week course, students will learn to implement, train and debug their own neural networks and gain a detailed understanding of cutting-edge research in computer vision. The final assignment will involve training a multi-million parameter convolutional neural network and applying it on the largest image classification dataset (ImageNet). We

fast.ai

Making neural nets
uncool again

[Home](#)

[About](#)

[Our MOOC](#)

[Posts by Topic](#)

© fast.ai 2018. All rights reserved.

Our courses (all are free and have no ads):

- Deep Learning Part 1: [Practical Deep Learning for Coders](#)
 - [Why we created the course](#)
 - [What we cover in the course](#)
- Deep Learning Part 2: [Cutting Edge Deep Learning for Coders](#)
- Computational Linear Algebra: [Online textbook](#) and [Videos](#)
- [Providing a Good Education in Deep Learning](#)—our teaching philosophy
- [A Unique Path to Deep Learning Expertise](#)—our teaching approach

fast.ai in the news:

- [20 Incredible Women Advancing A.I. Research](#)
- [AI Revolutionaries: Possible Futures](#)
- [The Digital Industrial Revolution](#)
- [Artificial Intelligence Education Transforms The Developing World](#)

A Discussion about Accessibility in AI at Stanford

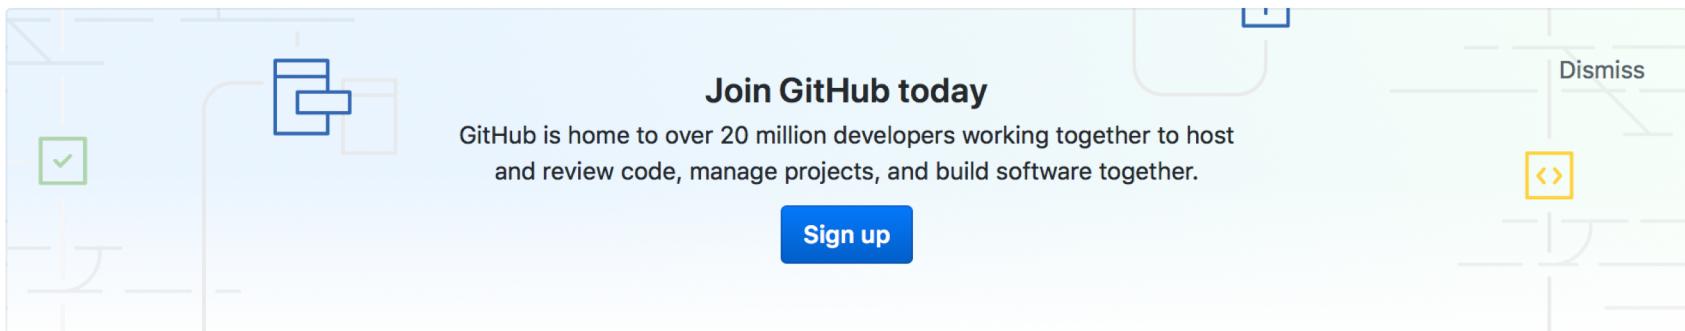
 Code

 Issues 22

 Pull requests 11

 Projects 0

 Insights



Jupyter notebooks for the code samples of the book "Deep Learning with Python"

 16 commits

 1 branch

 0 releases

 2 contributors

 MIT

Branch: master ▾

New pull request

Find file

Clone or download ▾

 fchollet committed on Nov 20, 2017 Merge pull request #11 from hiroyachiba/master ...		Latest commit 8a30b90 on Nov 20, 2017
2.1-a-first-look-at-a-neural-network.ipynb	Add notebooks	7 months ago
3.5-classifying-movie-reviews.ipynb	Add notebooks	7 months ago
3.6-classifying-newswires.ipynb	Add notebooks	7 months ago
3.7-predicting-house-prices.ipynb	typos	7 months ago
4.4-overfitting-and-underfitting.ipynb	typos	7 months ago

Deep Learning

An MIT Press book

Ian Goodfellow and Yoshua Bengio and Aaron Courville

[Exercises](#) [Lectures](#) [External Links](#)

The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free.

The deep learning textbook can now be ordered on [Amazon](#).

For up to date announcements, join our [mailing list](#).

Citing the book

To cite this book, please use this bibtex entry:

```
@book{Goodfellow-et-al-2016,
  title={Deep Learning},
  author={Ian Goodfellow and Yoshua Bengio and Aaron Courville},
  publisher={MIT Press},
  note={\url{http://www.deeplearningbook.org}},
  year={2016}
}
```

To write your own document using our LaTeX style, math notation, or to copy our notation page, download our [template](#) files.

[Errata in published editions](#)

VIBOT 2018

Closing remarks

Francesco Ciompi

francesco.ciompi@radboudumc.nl